

Exercise 12.15. It's-a me, Mario! This problem is inspired by Super Mario World, where Mario must make it from some starting point to the flag in front of a castle, collecting as many coins as possible along the way.

Let $G = (V, E)$ be a directed graph where each vertex $v \in V$ has $C_v \geq 0$ coins. For a walk in G , we say that the *number of coins collected by the walk* is the total sum of C_v over all *distinct* vertices v in the walk. (If we visit a vertex v more than once, we still only get C_v coins total.) Given $s, t \in V$, the goal is to compute the maximum number of coins collected by any (s, t) -walk.

Exercise 12.15.1. Let G be a DAG. For this problem, either (a) design and analyze a polynomial time algorithm (the faster the better), or (b) prove that a polynomial time algorithm would imply a polynomial time algorithm for SAT.

Solution. We claim that the following algorithm suffices.

mcw-dag(v, t):

/ given $v, t \in V$ in a DAG, computes the maximum number of coins collected by any (v, t) -walk and returns $-\infty$ if no such walk exists. */*

1. If $v = t$ then return C_v .
2. Let $m \leftarrow -\infty$.
3. For $(v, w) \in \delta^+(v)$ do
 - A. $m \leftarrow \max(m, C_v + \text{mcw-dag}(w, t))$
4. Return m

We find the maximum number of coins collected along any (s, t) -walk by calling mcw-dag(s, t).

Runtime. Given a goal vertex $t \in V$, if we cache the return value of mcw-dag(v, t) for all $v \in V$, then our algorithm has the following runtime complexity:

$$O\left(\sum_{v \in V} (1 + d^+(v))\right) = O(m + n)$$

■

Correctness. The correctness of this algorithm follows from performing induction, in reverse topological order, according to the recursive specification. We claim that mcw-dag(v, t) returns the maximum number of coins collected by any (v, t) -walk for all $v \in V$.

In the base case, v is the last vertex in topological order—a sink—so the maximum number of collectible coins is C_v if $v = t$ and $-\infty$ if $v \neq t$, as returned in the algorithm.

Now we assume that for some $v \in V$, the claim holds for all vertices that follow v in topological order. In the case where v is a sink, there are no vertices reachable from v and consequently, the algorithm correctly returns C_v if $v = t$ and $-\infty$ if $v \neq t$. Otherwise, since the claim holds, by assumption, for all w such that $(v, w) \in \delta^+(v)$, we can conclude that the maximum number of collectible coins on a (v, t) -walk is

$$C_v + \max_{w: (v, w) \in \delta^+(v)} \text{mcw-dag}(w, t)$$

as computed in the algorithm.

Thus, by induction, the claim that mcw-dag(v, t) returns the maximum collectible coins along a (v, t) -walk holds for all $v \in V$. ■

□

Exercise 12.15.2. Let G be a general directed graph. For this problem, either (a) design and analyze a polynomial time algorithm (the faster the better), or (b) prove that a polynomial time algorithm would imply a polynomial time algorithm for SAT.

Solution. When G is a general directed graph, we compress G to its condensation graph G' , then apply the algorithm given in 12.15.1 to G' .

mcw(s, t):

/ given $s, t \in V$ in a general digraph, computes the maximum number of coins collected by any (s, t) -walk and returns $-\infty$ if no such walk exists. */*

/ see 12.15.1 for mcw-dag() pseudocode. */*

1. Let $\{S_1, S_2, \dots, S_k\}$ be the strongly-connected components in G .
2. Contract each S_i into a single vertex s_i in $G' = (V', E')$ with $C_{S_i} = \sum_{v \in S_i} C_v$.
3. For $i \neq j$, $(s_i, s_j) \in E'$ iff there exist $u \in S_i, v \in S_j$ with $(u, v) \in E$.
4. Let S_s and S_t be the sccs containing s and t , respectively.
5. Return mcw-dag(s_s, s_t).

Calling mcw(s, t) gives us the maximum number of coins collected by any (s, t) -walk.

Runtime. The runtime of SCC-compression on G is in $O(m + n)$, as is the runtime of mcw-dag(). Hence, the total runtime complexity of our algorithm is in $O(m + n)$. ■

Correctness. Since the correctness of mcw-dag() was proven in 12.15.1, it suffices for us to show that our SCC-compression preserves the maximal number of coins collected along any (s, t) -walk.

First, we claim that a maximal coin-collection walk in G necessarily collects all of the coins in every SCC it visits; otherwise, we would be able to make “detours” in at least one SCC to pick up coins we’ve missed, contradicting the maximality of the walk. We call such a walk an SCC-walk.

Since we now know the maximum number of coins collected from an (s, t) -walk in G can be given by an SCC-walk on G ; so, we can simply maximize coin collection on only the SCC-walks in G , rather than on all walks.

We have a clear correspondence between coins collected on SCC-walks in G and all walks in G' :

- (i) First, for every (s, t) -SCC-walk W on G , there exists a walk of the same coin-numerage from s_s to s_t in G' by taking all the SCCs W passes through.
- (ii) Conversely, for every walk from s_s to s_t in G' , there exists an (s, t) -SCC-walk in G with the same total weight.

By this bijection, we can conclude that the maximum coin collection over all (s, t) -walks in G is equal to the maximum coin collection over all (s_s, s_t) -walks in G' . ■

□